

analysis

October 10, 2023

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sn
import numpy as np

[3]: files = ["asian", "black", "indian", "latino_hispanic", "maori", "middle_eastern", "white", "pasifika"]

attributes = ["Gender", "Ethnicity", "Mood", "Age"]

def get_range(age_range):
    if age_range == "more than 70":
        return range(70, 120)
    bounds = str(age_range).split("-")

    if len(bounds) == 1:
        lower = (int(int(bounds[0]) / 10)) * 10
        return range(lower, lower + 10)

    return range(int(bounds[0]), int(bounds[1]) + 1)

def get_str_range(age):
    if int(age) >= 70:
        return "more than 70"
    r = get_range(age)
    return f"{r[0]}--{r[9]}"

def get_accuracies(df):
    c_results = {}
    for attr in attributes:
        gb = df.groupby(f"{attr} accurate")[f"{attr} accurate"].count()
        c_results[attr] = [(gb.index[i], gb[i]) for i in range(len(gb.index))]

    return c_results

[4]: dfs = {}

per_category_results = {}
```

```

for f in files:
    path = f"{f}_stats.csv"
    df = pd.read_csv(path, index_col = "Unnamed: 0")
    df["True Ethnicity"] = df["True Ethnicity"].str.replace("_", " ").str.
    ↪lower()
    df["Classified Mood"] = df["Classified Mood"].str.replace("surprise", "
    ↪surprised")
    df["Classified Mood"] = df["Classified Mood"].str.replace("disgust", "
    ↪disgusted")
    if f == "pasifika" or f == "maori":
        df = df.sample(n = 100)
        df["True Gender"] = df["True Gender"].str.replace("Female", "Woman")
        df["True Gender"] = df["True Gender"].str.replace("Male", "Man")
        df["True Ethnicity"] = df["True Ethnicity"].str.replace("european", "
    ↪maori")
        df["True Age"] = df["True Age"].apply(get_str_range)

    df.insert(3, "Gender accurate", df["True Gender"] == df["Classified
    ↪Gender"])
    df.insert(7, "Ethnicity accurate", df["True Ethnicity"] == df["Classified
    ↪Ethnicity"])
    df.insert(11, "Mood accurate", df["True Mood"] == df["Classified Mood"])
    df.insert(15, "Age accurate", df.apply(lambda row: row["Classified Age"] in
    ↪get_range(row["True Age"]), axis = 1))
    accuracy = get_accuracies(df)

    per_category_results[f] = accuracy

    dfs[f] = df

```

```

[5]: # All categories
all_df = pd.concat(dfs.values())

all_acc = get_accuracies(all_df)

```

```

[6]: # Matrix generation

def confusion_matrix(group, attr, options, normalise = True):
    row = {}
    correct_option = list(group[f"True {attr}"])[0]
    for o in options:
        correct = correct_option == o
        coeff = (100 * (1 / len(group)) * (1 if correct else -1)) if normalise
    ↪else 1
        if attr == "Age":
            correct = correct_option == get_str_range(o)

```

```

        coeff *= -correct if correct and normalise else 1
        num_correct = group[["True Age", "Classified Age"]]\
            .apply(lambda row: (row["Classified Age"]
        ↪Age"] in get_range(row["True Age"])) & (get_str_range(row["Classified Age"])\
        ↪== get_str_range(o)), axis = 1)\
            .sum()

        total_classified = coeff * group[["True Age", "Classified Age"]]\
            .apply(lambda row:
        ↪int(row["Classified Age"] in get_range(o)), axis = 1)\
            .sum()

        row[get_str_range(o)] = (num_correct, total_classified)

    else:
        num_correct = sum((group[f"Classified {attr}"] == group[f"True
        ↪{attr}"]) & (group[f"Classified {attr}"] == o))
        total_classified = coeff * sum(group[f"Classified {attr}"] == o)
        row[o] = (num_correct, total_classified)

    return pd.Series(row)

norm_matrices = {}
for attr in attributes:
    options = set(all_df[f"Classified {attr}"])
    matrix = all_df.groupby(f"True {attr}").apply(confusion_matrix, attr =
    ↪attr, options = options, normalise = True)
    new_order = list(filter(lambda i: i in matrix.columns, matrix.index))
    new_order = new_order + list(filter(lambda c: c not in new_order, matrix.
    ↪columns))
    matrix = matrix[list(new_order)]
    norm_matrices[attr] = matrix.applymap(lambda cell: cell[1])

c_matrices = {}
for att1 in attributes:
    c_matrices[att1] = {}
    for att2 in attributes:
        options = set(all_df[f"Classified {att2}"])
        matrix = all_df.groupby(f"True {att1}").apply(confusion_matrix, attr =
        ↪att2, options = options, normalise = False)
        c_matrices[att1][att2] = matrix

```

[13]: # Precision/recall calculations

```

pr_results = {}
for attr, matrix in norm_matrices.items():

```

```

prs = []
for c in matrix.columns:
    tp = matrix[c].loc[[c]][0]
    fp = sum(matrix[c].abs()) - tp
    fn = matrix.loc[[c]].abs().sum(axis = 1)[0] - tp
    tn = all_df.groupby(f"True {attr}").size().sum() - all_df.
    ↳groupby(f"True {attr}").size()[c]
    precision = tp / (tp + fp)
    recall = tp / (tp + fn)
    specificity = tn / (tn + fp)

    prs.append((c, precision, recall, specificity))
pr_results[attr] = prs

```

1 Visualisation

1.1 Accuracies

1.1.1 Total accuracy

```

[8]: ethnicities = files

fig, ax = plt.subplots(layout = "constrained")
fig.set_size_inches(9, 5)
width = 0.3
x = np.arange(len(attributes))

false_heights = [all_acc[attr][0][1] / len(ethnicities) for attr in attributes]
true_heights = [100 - h for h in false_heights]
print(false_heights)

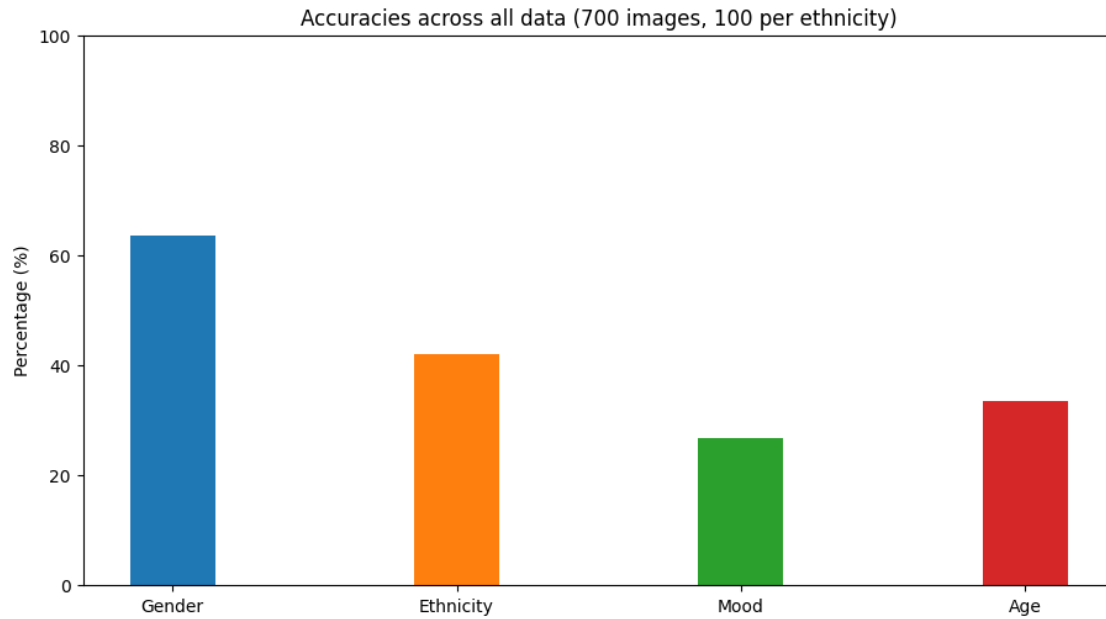
ax.bar(x, true_heights, width, color = ["tab:blue", "tab:orange", "tab:green",
    ↳"tab:red"])

ax.set_ylabel('Percentage (%)')
ax.set_title('Accuracies across all data (700 images, 100 per ethnicity)')
ax.set_xticks(x, attributes)
ax.set_ylim(0, 100)

```

[36.5, 58.0, 73.25, 66.5]

[8]: (0.0, 100.0)



1.1.2 Accuracy by ethnicity

```
[9]: fig, ax = plt.subplots(layout = "constrained")
fig.set_size_inches(9, 5)
multiplier = 0
width = 0.15
x = np.arange(len(ethnicities))

for attr in attributes:
    false_heights = [per_category_results[d][attr][0][1] for d in
↳per_category_results]
    true_heights = [100 - x for x in false_heights]

    offset = width * multiplier
    ax.bar(x + offset, true_heights, width, label = attr)

    multiplier += 1

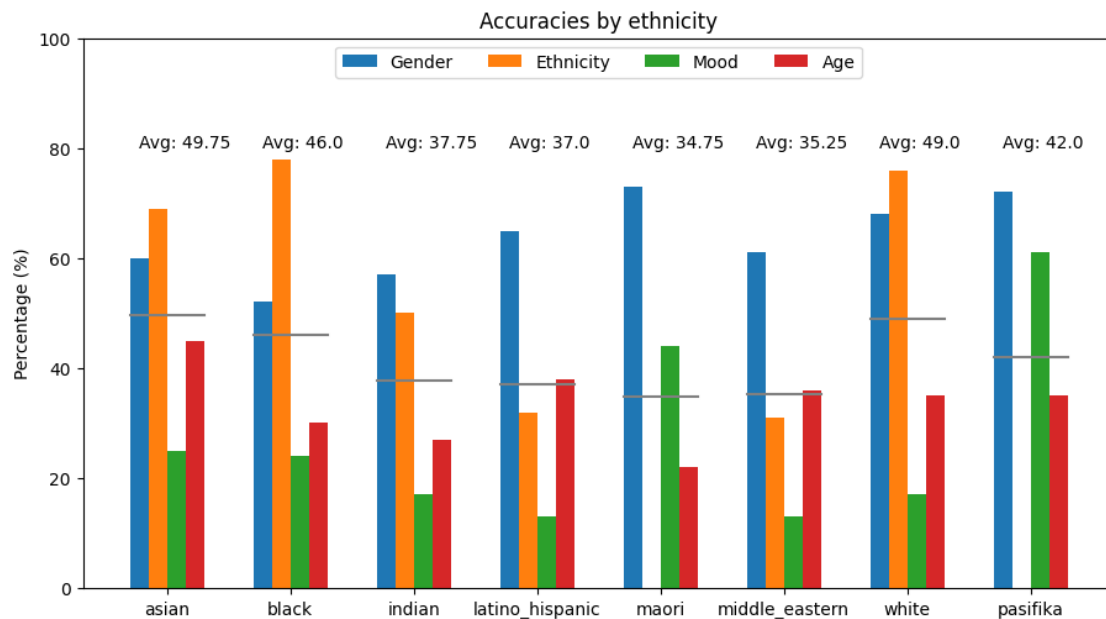
for eth in range(len(ethnicities)):
    avg = sum(100 - per_category_results[ethnicities[eth]][x][0][1] for x in
↳per_category_results[ethnicities[eth]]) / len(attributes)
    ax.text(x[eth], 80, f"Avg: {avg}")
    ax.plot([x[eth] - width / 2, x[eth] + width * 3.5], [avg, avg], color =
↳"gray")
```

```

ax.set_ylabel('Percentage (%)')
ax.set_title('Accuracies by ethnicity')
ax.set_xticks(x + 1.5*width, ethnicities)
ax.legend(loc='upper center', ncols = 4)
ax.set_ylim(0, 100)

```

[9]: (0.0, 100.0)



1.1.3 Accuracy matrices

```

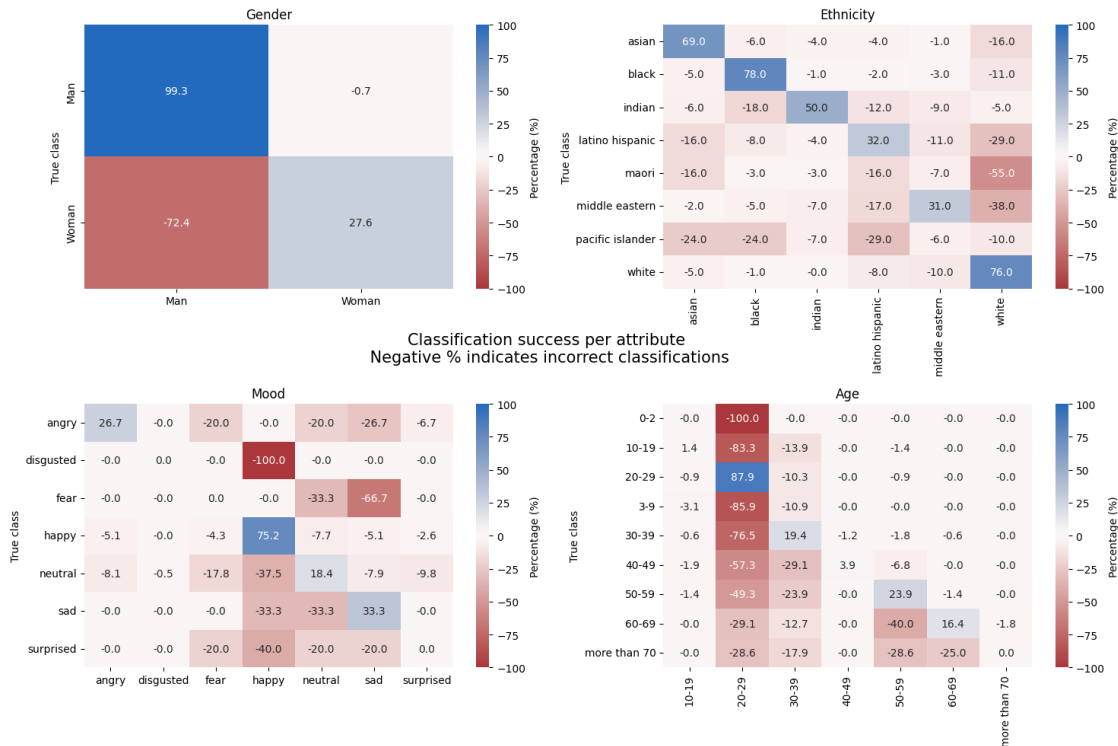
[10]: fig, axes = plt.subplots(2, 2, figsize = (15, 10))

for i, (key, data) in enumerate(norm_matrices.items()):
    row, col = divmod(i, 2)
    sn.heatmap(data, fmt = '.1f', annot = True, cmap = "vlag_r", ax = axes[row, col],
               vmin = -100, vmax = 100,
               cbar_kws = {"label": "Percentage (%)"})
    axes[row, col].set_title(key)
    axes[row, col].set_ylabel("True class")

plt.tight_layout()
plt.figtext(0.48, 0.52, "Classification success per attribute\n Negative %\n indicates incorrect classifications", ha = "center", fontsize = 15)

```

[10]: Text(0.48, 0.52, 'Classification success per attribute\n Negative % indicates incorrect classifications')



1.2 Precision, Recall, Sensitivity and Specifity


```
[14]: fig, axes = plt.subplots(2, 2, figsize = (15, 10))
width = 0.2
for i, (key, data) in enumerate(pr_results.items()):
    row, col = divmod(i, 2)
    multiplier = 0
    x = np.arange(len(data))

    for i in range(3):
        offset = width * multiplier
        axes[row, col].bar(x + offset, [t[i + 1] for t in data], width, label =
        ["Precision", "Recall", "Specifity"][i])

        multiplier += 1

    axes[row, col].set_title(key)
    axes[row, col].set_xticks(x + width, [t[0] for t in data])
    axes[row, col].set_ylim(0, 1)
    axes[row, col].legend(loc = 'upper center', ncols = 3)
```

```
plt.tight_layout()
```





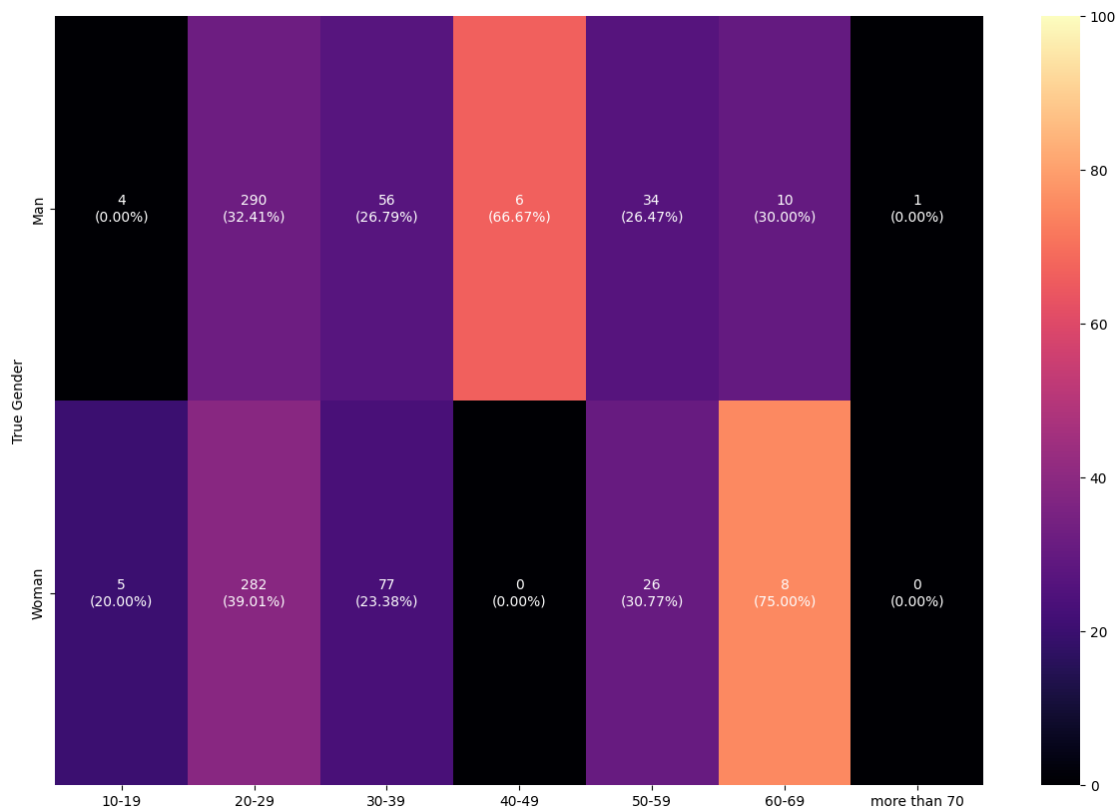
1.3 Bias heatmaps

```
[12]: count = 0
for (att1, inner) in c_matrices.items():
    for (att2, matrix) in inner.items():
        if att1 == att2:
            continue
        figure = plt.figure(figsize = (15, 10))
        annot_matrix = matrix.applymap(lambda t: str(t[1]) + f"\n({(100 * t[0] /
↪t[1]) if t[1] != 0 else 0:.2f}%)")
        sn.heatmap(matrix.applymap(lambda t: (100 * t[0] / t[1]) if t[1] != 0
↪else 0),
                    vmin = 0,
                    vmax = 100,
                    cmap = "magma",
                    fmt = 's',
                    annot = annot_matrix)
        plt.show()

plt.tight_layout()
```









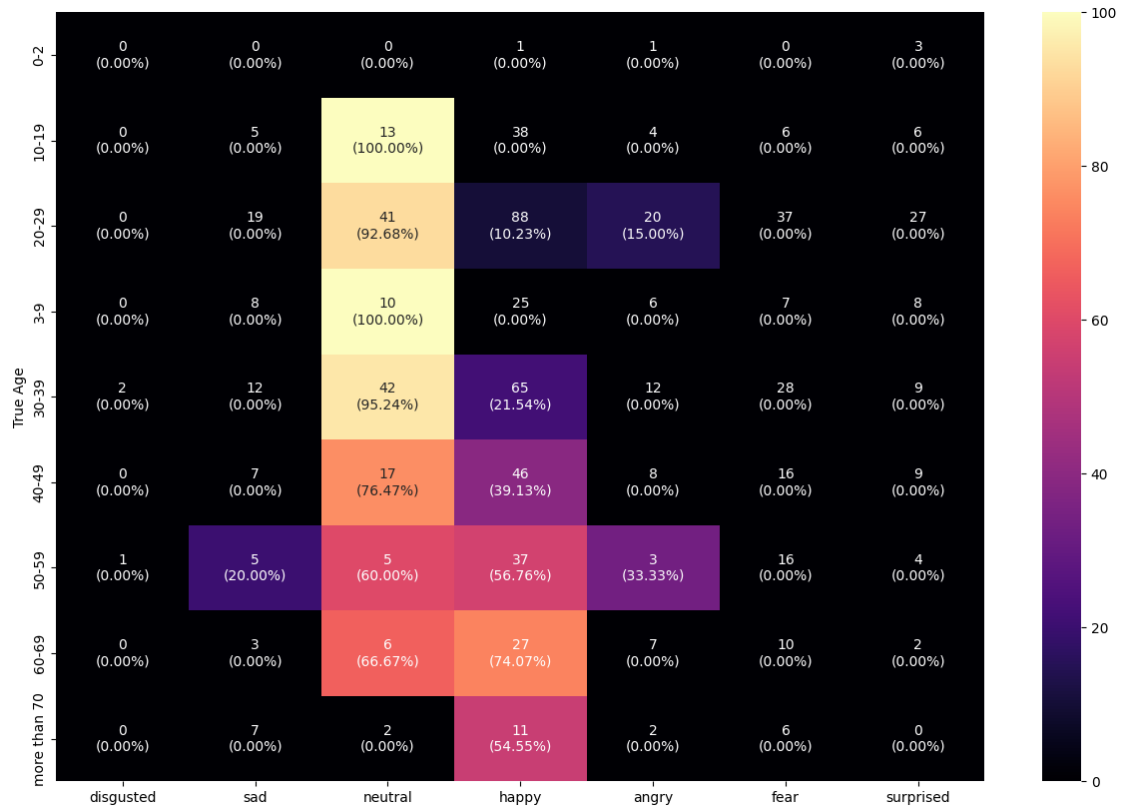












<Figure size 640x480 with 0 Axes>